

Model Development Phase

Date	09 - JULY- 2024
Team ID	SWTID1719999219
Project Title	Crystal Clear Vision: Revolutionizing Cataract Prediction through Transfer Learning Mastery
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include a summary and training and validation performance metrics for multiple models, presented through respective screenshots.

Initial Model Training Code (5 marks):

Paste the screenshot of the model training code

RESNET50 -

Resnet50

First, we will do feature extraction with keras functional API

```
[ ] #We are going to use EarlyStopping. Basically what it does is, it monitors validation accuracy.
    #If validation accuracy does not improve for 3 epochs, the model fitting will cease.
    from keras.callbacks import EarlyStopping, ModelCheckpoint
```

[illegible]

VGG16 –

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

# Load VGG16 base model
vgg = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze VGG16 layers
for layer in vgg.layers:
    layer.trainable = False

# Create your model on top of VGG16
x = Flatten()(vgg.output)
output = Dense(1, activation='sigmoid')(x)
model = Model(vgg.input, output)

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])
```

```
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
train_path = os.path.join(base_path, "train")
validation_path = os.path.join(base_path, "validation")
model_save_path = os.path.join(base_path, "model/vgg16_model.h5")

train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
validation_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(train_path, target_size=(224, 224), batch_size=16, class_mode='binary')
validation_generator = validation_datagen.flow_from_directory(validation_path, target_size=(224, 224), batch_size=16, class_mode='binary')

##model.fit(train_generator, validation_data=validation_generator, epochs=2)
# Define callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
checkpoint = ModelCheckpoint(model_save_path, monitor='val_loss', save_best_only=True, verbose=1)

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])

# Train the model with callbacks
history = model.fit(
    train_generator,
    validation_data=validation_generator,
    epochs=50,
    callbacks=[early_stopping, checkpoint]
)

# Print the model save path
print(f"Model saved at: {model_save_path}")
```

```
base_model= tf.keras.applications.EfficientNetB1(include_top=False)
#include top is false as we are going to create our own output layers for the model

base_model.trainable=False
#freezing the base layers

#input layer
inputs=tf.keras.layers.Input(shape=(240,240,3),name='input_layer')

#adding in data augmentation as a layer itself
x=data_augmentation(inputs)

#pass inputs to base model
x=base_model(x, training=False)
print(x.shape)

#We will use average pooling to reduce the size of the feature map.
x=tf.keras.layers.GlobalAveragePooling2D(name='global_average_pooling_layer')(x)
print(x.shape)

outputs=tf.keras.layers.Dense(1,activation='sigmoid',name='output_layer')(x)
model_0= tf.keras.Model(inputs,outputs)
```

```
#define early stopping callback.
early_stopping=EarlyStopping(monitor='val_accuracy',patience=3,restore_best_weights=True)
checkpoint = ModelCheckpoint('best_efficientnet_model.h5', monitor='val_accuracy', save_best_only=True)

model_0.compile(loss='binary_crossentropy',
                optimizer=tf.keras.optimizers.Adam(),
                metrics=['accuracy'])

history_model_0= model_0.fit(train_data,
                             epochs=50,
                             validation_data=val_data,
                             callbacks=[early_stopping,checkpoint])
```

InceptionV3 Model –

```
# InceptionV3 model
inception = InceptionV3(include_top=False, input_shape=(299, 299, 3))
x = Flatten()(inception.output)
output = Dense(2, activation='softmax')(x) # Changed to 2 classes for binary classification

inception_model = Model(inception.input, output)
inception_model.summary()

# Compile the model
inception_model.compile(loss='categorical_crossentropy', optimizer=Adam(learning_rate=0.001), metrics=['accuracy'])

# Define Early Stopping and Model Checkpoint callbacks
early_stopping = EarlyStopping(monitor='val_accuracy', patience=3, restore_best_weights=True)
checkpoint = ModelCheckpoint('best_inception_model.h5', monitor='val_accuracy', save_best_only=True)

# Train the model with early stopping and model checkpoint callbacks
history = inception_model.fit(train_gen, validation_data=val_gen, epochs=epochs, callbacks=[early_stopping, checkpoint])
```

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics																																							
Resnet50	<pre>model_0.summary()</pre> <p>Model: "model"</p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>input layer (InputLayer)</td><td>[(None, 224, 224, 3)]</td><td>0</td></tr> <tr> <td>data_augmentation (Sequential)</td><td>(None, None, None, 3)</td><td>0</td></tr> <tr> <td>resnet50 (Functional)</td><td>(None, None, None, 2048)</td><td>23587712</td></tr> <tr> <td>global_average_pooling_layer (GlobalAveragePooling2D)</td><td>(None, 2048)</td><td>0</td></tr> <tr> <td>output_layer (Dense)</td><td>(None, 1)</td><td>2049</td></tr> </table> <p> Total params: 23589761 (89.99 MB) Trainable params: 2049 (8.00 KB) Non-trainable params: 23587712 (89.98 MB) </p>	Layer (type)	Output Shape	Param #	input layer (InputLayer)	[(None, 224, 224, 3)]	0	data_augmentation (Sequential)	(None, None, None, 3)	0	resnet50 (Functional)	(None, None, None, 2048)	23587712	global_average_pooling_layer (GlobalAveragePooling2D)	(None, 2048)	0	output_layer (Dense)	(None, 1)	2049	<pre> Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5 160/160 [=====] 6s 0.0s/step Epoch 1/50 10/10 [=====] - ETA: 0s - loss: 0.5202 - accuracy: 0.7625/val_loss: 0.4108 - val_accuracy: 0.8000 10/10 [=====] - saving_model 10/10 [=====] - 31s 24s/step - loss: 0.5092 - accuracy: 0.7625 - val_loss: 0.4108 - val_accuracy: 0.8000 10/10 [=====] - 36s 24s/step - loss: 0.4887 - accuracy: 0.7900 - val_loss: 0.3812 - val_accuracy: 0.8000 10/10 [=====] - 34s 24s/step - loss: 0.4099 - accuracy: 0.8406 - val_loss: 0.3441 - val_accuracy: 0.8500 Epoch 4/50 10/10 [=====] - 37s 24s/step - loss: 0.3131 - accuracy: 0.8719 - val_loss: 0.3630 - val_accuracy: 0.8750 Epoch 5/50 10/10 [=====] - 31s 24s/step - loss: 0.2973 - accuracy: 0.8750 - val_loss: 0.3491 - val_accuracy: 0.8500 Epoch 6/50 10/10 [=====] - 34s 24s/step - loss: 0.2981 - accuracy: 0.8750 - val_loss: 0.3415 - val_accuracy: 0.8500 Epoch 7/50 10/10 [=====] - 29s 11s/step - loss: 0.2885 - accuracy: 0.8906 - val_loss: 0.3081 - val_accuracy: 0.8750 </pre>																					
Layer (type)	Output Shape	Param #																																							
input layer (InputLayer)	[(None, 224, 224, 3)]	0																																							
data_augmentation (Sequential)	(None, None, None, 3)	0																																							
resnet50 (Functional)	(None, None, None, 2048)	23587712																																							
global_average_pooling_layer (GlobalAveragePooling2D)	(None, 2048)	0																																							
output_layer (Dense)	(None, 1)	2049																																							
VGG16	<pre># Print the model summary model.summary()</pre> <p> Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 160/160 [=====] 6s 0.0s/step Model: "model" </p> <table> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> <tr> <td>input_1 (InputLayer)</td><td>[(None, 224, 224, 3)]</td><td>0</td></tr> <tr> <td>block1_conv1 (Conv2D)</td><td>(None, 224, 224, 64)</td><td>1792</td></tr> <tr> <td>block1_conv2 (Conv2D)</td><td>(None, 224, 224, 64)</td><td>36928</td></tr> <tr> <td>block1_pool (MaxPooling2D)</td><td>(None, 112, 112, 64)</td><td>0</td></tr> <tr> <td>block2_conv1 (Conv2D)</td><td>(None, 112, 112, 128)</td><td>73856</td></tr> <tr> <td>block2_conv2 (Conv2D)</td><td>(None, 112, 112, 128)</td><td>147584</td></tr> <tr> <td>block2_pool (MaxPooling2D)</td><td>(None, 56, 56, 128)</td><td>0</td></tr> <tr> <td>block3_conv1 (Conv2D)</td><td>(None, 56, 56, 256)</td><td>295168</td></tr> <tr> <td>block3_conv2 (Conv2D)</td><td>(None, 56, 56, 256)</td><td>590080</td></tr> <tr> <td>block3_conv3 (Conv2D)</td><td>(None, 56, 56, 256)</td><td>590080</td></tr> <tr> <td>block3_pool (MaxPooling2D)</td><td>(None, 28, 28, 256)</td><td>0</td></tr> <tr> <td>block4_conv1 (Conv2D)</td><td>(None, 28, 28, 512)</td><td>1180160</td></tr> </table>	Layer (type)	Output Shape	Param #	input_1 (InputLayer)	[(None, 224, 224, 3)]	0	block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792	block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928	block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0	block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856	block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584	block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0	block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168	block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080	block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080	block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0	block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160	<pre> Found 120 images belonging to 2 classes. Found 40 images belonging to 2 classes. Epoch 1/50 20/20 [=====] - ETA: 0s - loss: 0.5035 - accuracy: 0.7906 Epoch 1: val_loss improved from inf to 0.38221, saving model to /content/repository/yliachen84-retina_dataset-9140ef4/model/vgg16_model.h5 20/20 [=====] - 215s 10s/step - loss: 0.5035 - accuracy: 0.7906 - val_loss: 0.3822 - val_accuracy: 0.8500 /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:1310: UserWarning: You are saving your model as an H5 file via `model.save('saving_vgg16_model.h5')`. Epoch 1: val_loss did not improve from 0.38221 Epoch 2: val_loss did not improve from 0.38221 20/20 [=====] - 215s 11s/step - loss: 0.2866 - accuracy: 0.8687 - val_loss: 0.4177 - val_accuracy: 0.8500 Epoch 3/50 20/20 [=====] - ETA: 0s - loss: 0.2477 - accuracy: 0.8938 Epoch 3: val_loss improved from 0.38221 to 0.35826, saving model to /content/repository/yliachen84-retina_dataset-9140ef4/model/vgg16_model.h5 20/20 [=====] - 205s 11s/step - loss: 0.2477 - accuracy: 0.8938 - val_loss: 0.3582 - val_accuracy: 0.8500 Epoch 4/50 20/20 [=====] - ETA: 0s - loss: 0.2544 - accuracy: 0.8781 Epoch 4: val_loss did not improve from 0.35826 20/20 [=====] - ETA: 0s - loss: 0.2544 - accuracy: 0.8781 Epoch 4: val_loss did not improve from 0.35826 Epoch 5/50 20/20 [=====] - 205s 11s/step - loss: 0.2544 - accuracy: 0.8781 - val_loss: 0.3826 - val_accuracy: 0.8500 Epoch 6/50 20/20 [=====] - ETA: 0s - loss: 0.2477 - accuracy: 0.8969 Epoch 6: val_loss did not improve from 0.35826 Epoch 7/50 20/20 [=====] - 211s 10s/step - loss: 0.2208 - accuracy: 0.8969 - val_loss: 0.4516 - val_accuracy: 0.8500 Epoch 8/50 20/20 [=====] - ETA: 0s - loss: 0.2097 - accuracy: 0.9094 Epoch 8: val_loss did not improve from 0.35826 20/20 [=====] - 209s 10s/step - loss: 0.2097 - accuracy: 0.9094 - val_loss: 0.3616 - val_accuracy: 0.9000 Epoch 9/50 20/20 [=====] - ETA: 0s - loss: 0.2100 - accuracy: 0.9094 Epoch 9: val_loss did not improve from 0.35826 Epoch 10/50 20/20 [=====] - 208s 10s/step - loss: 0.2160 - accuracy: 0.9094 - val_loss: 0.3668 - val_accuracy: 0.7500 Epoch 11/50 20/20 [=====] - ETA: 0s - loss: 0.2170 - accuracy: 0.9062 Epoch 11: val_loss did not improve from 0.35826 Epoch 12/50 20/20 [=====] - 214s 11s/step - loss: 0.2170 - accuracy: 0.9062 - val_loss: 0.4116 - val_accuracy: 0.8500 Model saved at: /content/repository/yliachen84-retina_dataset-9140ef4/model/vgg16_model.h5 </pre>
Layer (type)	Output Shape	Param #																																							
input_1 (InputLayer)	[(None, 224, 224, 3)]	0																																							
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792																																							
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928																																							
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0																																							
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856																																							
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584																																							
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0																																							
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168																																							
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080																																							
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080																																							
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0																																							
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160																																							

	<pre> block4_conv2 (Conv2D) (None, 28, 28, 512) 2359808 block4_conv3 (Conv2D) (None, 28, 28, 512) 2359808 block4_pool (MaxPooling2D) (None, 14, 14, 512) 0 block5_conv1 (Conv2D) (None, 14, 14, 512) 2359808 block5_conv2 (Conv2D) (None, 14, 14, 512) 2359808 block5_conv3 (Conv2D) (None, 14, 14, 512) 2359808 block5_pool (MaxPooling2D) (None, 7, 7, 512) 0 flatten (Flatten) (None, 25088) 0 dense (Dense) (None, 1) 25089 ===== Total params: 14739777 (56.23 MB) Trainable params: 25089 (98.00 KB) Non-trainable params: 14714688 (56.13 MB) </pre>																			
EfficientnetB1	<pre> model_0.summary() Model: "model_1" </pre> <table> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>input_layer (InputLayer)</td><td>[(None, 240, 240, 3)]</td><td>0</td></tr> <tr> <td>data_augmentation (Sequential)</td><td>(None, None, None, 3)</td><td>0</td></tr> <tr> <td>efficientnetb1 (Functional)</td><td>(None, None, None, 1280)</td><td>6575239</td></tr> <tr> <td>global_average_pooling_layer (GlobalAveragePooling2D)</td><td>(None, 1280)</td><td>0</td></tr> <tr> <td>output_layer (Dense)</td><td>(None, 1)</td><td>1281</td></tr> </tbody> </table> <pre> ===== Total params: 6576520 (25.09 MB) Trainable params: 1281 (5.00 KB) Non-trainable params: 6575239 (25.08 MB) </pre>	Layer (type)	Output Shape	Param #	input_layer (InputLayer)	[(None, 240, 240, 3)]	0	data_augmentation (Sequential)	(None, None, None, 3)	0	efficientnetb1 (Functional)	(None, None, None, 1280)	6575239	global_average_pooling_layer (GlobalAveragePooling2D)	(None, 1280)	0	output_layer (Dense)	(None, 1)	1281	<pre> Epoch 1/50 10/10 [=====] - ETA: 0s - loss: 0.5813 - accuracy: 0.7031/usr/local/lib/python3.10/dist-packages/keras/src/saving_api_save_model(10/10 [=====] - 49% 2s/step - loss: 0.5833 - accuracy: 0.7031 - val_loss: 0.5492 - val_accuracy: 0.7500 Epoch 2/50 10/10 [=====] - 29% 1s/step - loss: 0.4924 - accuracy: 0.7594 - val_loss: 0.4763 - val_accuracy: 0.7500 Epoch 3/50 10/10 [=====] - 31% 1s/step - loss: 0.4256 - accuracy: 0.8094 - val_loss: 0.4104 - val_accuracy: 0.8000 Epoch 4/50 10/10 [=====] - 28% 1s/step - loss: 0.3824 - accuracy: 0.8562 - val_loss: 0.4005 - val_accuracy: 0.8000 Epoch 5/50 10/10 [=====] - 27% 1s/step - loss: 0.3291 - accuracy: 0.8594 - val_loss: 0.3827 - val_accuracy: 0.8000 Epoch 6/50 10/10 [=====] - 28% 1s/step - loss: 0.3217 - accuracy: 0.8719 - val_loss: 0.3631 - val_accuracy: 0.8250 Epoch 7/50 10/10 [=====] - 27% 930ms/step - loss: 0.3006 - accuracy: 0.8813 - val_loss: 0.3517 - val_accuracy: 0.8250 Epoch 8/50 10/10 [=====] - 26% 1s/step - loss: 0.2921 - accuracy: 0.8906 - val_loss: 0.3606 - val_accuracy: 0.8250 Epoch 9/50 10/10 [=====] - 26% 1s/step - loss: 0.2749 - accuracy: 0.8875 - val_loss: 0.3477 - val_accuracy: 0.8250 </pre>
Layer (type)	Output Shape	Param #																		
input_layer (InputLayer)	[(None, 240, 240, 3)]	0																		
data_augmentation (Sequential)	(None, None, None, 3)	0																		
efficientnetb1 (Functional)	(None, None, None, 1280)	6575239																		
global_average_pooling_layer (GlobalAveragePooling2D)	(None, 1280)	0																		
output_layer (Dense)	(None, 1)	1281																		
Inception	<pre> activation_186 (Activation) (None, 8, 8, 384) 0 ["batch_normalization_186[0][0]"] batch_normalization_187 (Batch Normalization) (None, 8, 8, 384) 576 ["conv2d_187[0][0]"] activation_179 (Activation) (None, 8, 8, 320) 0 ["batch_normalization_179[0][0]"] mixed0_1 (Concatenate) (None, 8, 8, 768) 0 ["activation_181[0][0]", "activation_182[0][0]"] concatenate_3 (Concatenate) (None, 8, 8, 768) 0 ["activation_185[0][0]", "activation_186[0][0]"] activation_187 (Activation) (None, 8, 8, 384) 0 ["batch_normalization_187[0][0]"] mixed0 (Concatenate) (None, 8, 8, 2048) 0 ["activation_179[0][0]", "mixed0_1[0][0]", "concatenate_3[0][0]", "activation_187[0][0]"] flatten_1 (Flatten) (None, 131072) 0 ["mixed0[0][0]"] dense_1 (Dense) (None, 2) 262240 ["flatten_1[0][0]"] ===== Total params: 22064930 (84.17 MB) Trainable params: 22000498 (84.04 MB) Non-trainable params: 34832 (134.50 KB) </pre>	<pre> Epoch 1/50 10/10 [=====] - ETA: 0s - loss: 5.6772 - accuracy: 0.7219 /usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:1188: KeyboardInterrupt: saving_api_save_model(10/10 [=====] - 38% 16s/step - loss: 5.6772 - accuracy: 0.7219 - val_loss: 821.4074 - val_accuracy: 0.7500 Epoch 2/50 10/10 [=====] - 41% 41s/step - loss: 1.6558 - accuracy: 0.8281 - val_loss: 15207.5112 - val_accuracy: 0.7500 Epoch 3/50 10/10 [=====] - 42% 41s/step - loss: 0.8089 - accuracy: 0.8089 - val_loss: 191006.0000 - val_accuracy: 0.7500 Epoch 4/50 10/10 [=====] - 34% 16s/step - loss: 0.3196 - accuracy: 0.8056 - val_loss: 1680414.0000 - val_accuracy: 0.7500 Model saved at: /content/next_inception_model.h5 2/2 [=====] - 1% 2s/step - loss: 0.7500 - accuracy: 0.7500 Test accuracy: 0.75 2/2 [=====] - 1% 2s/step </pre>																		